



Automation fails at the boring part

Supporting line

The model was fine

A postmortem traced a double post to the retry logic, not the AI.

Three questions decide it

Who approves it? What if it runs twice? Where are the numbers checked?

Build the safety layer first

One owner per step, a human approval, and one home for the rules.

Trust is the product

People rely on a system they can predict.

Which step keeps a human in it?

Supporting line